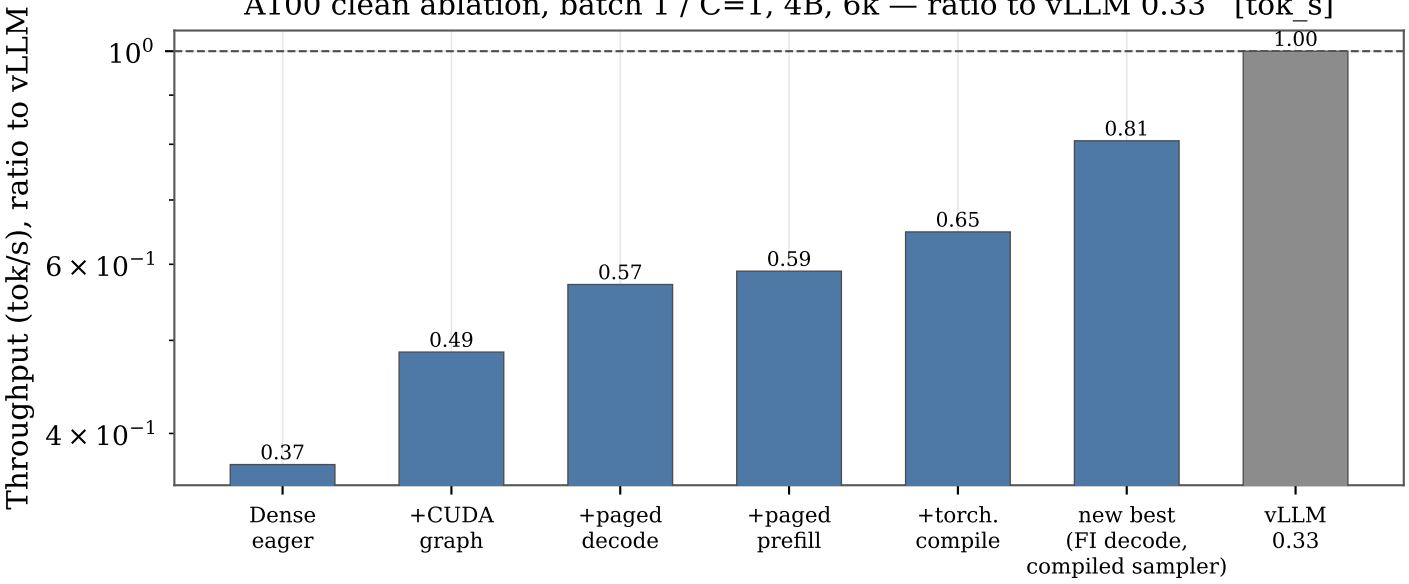
A100 clean ablation, batch 1 / C=1, 4B, 6k — ratio to vLLM 0.33 [tok_s]
Throughput (tok/s), ratio to vLLM
10^{0}
6×10^{-1}
4×10^{-1}
0.37
0.49
0.57
0.59
0.65
0.81
1.00
Dense eager
+CUDA graph
+paged decode
+paged prefill
+torch. compile
new best (FI decode, compiled sampler)
vLLM 0.33