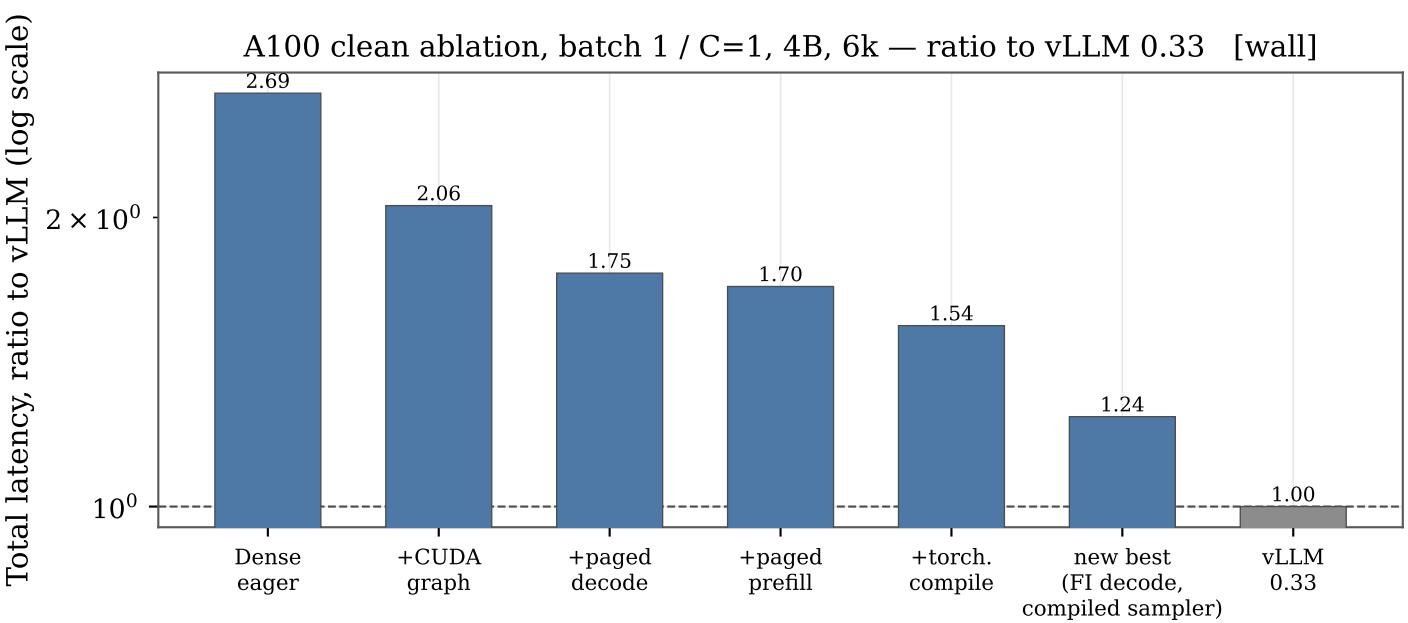
A100 clean ablation, batch 1 / C=1, 4B, 6k — ratio to vLLM 0.33 [wall]
Total latency, ratio to vLLM (log scale)
2×10^0
10^0
2.69
2.06
1.75
1.70
1.54
1.24
1.00
Dense eager
+CUDA graph
+paged decode
+paged prefill
+torch. compile
new best (FI decode, compiled sampler)
vLLM 0.33