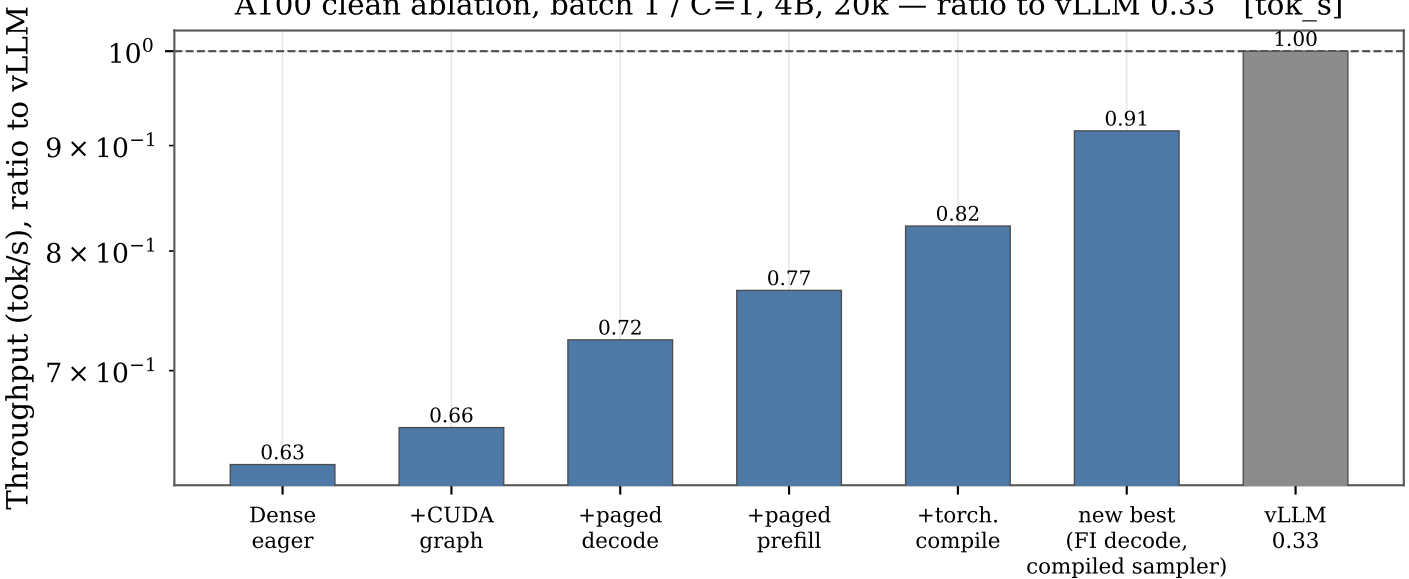
A100 clean ablation, batch 1 / C=1, 4B, 20k — ratio to vLLM 0.33 [tok_s]
Throughput (tok/s), ratio to vLLM
10^{0}
9×10^{-1}
8×10^{-1}
7×10^{-1}
0.63
0.66
0.72
0.77
0.82
0.91
1.00
Dense eager
+CUDA graph
+paged decode
+paged prefill
+torch. compile
new best (FI decode, compiled sampler)
vLLM 0.33