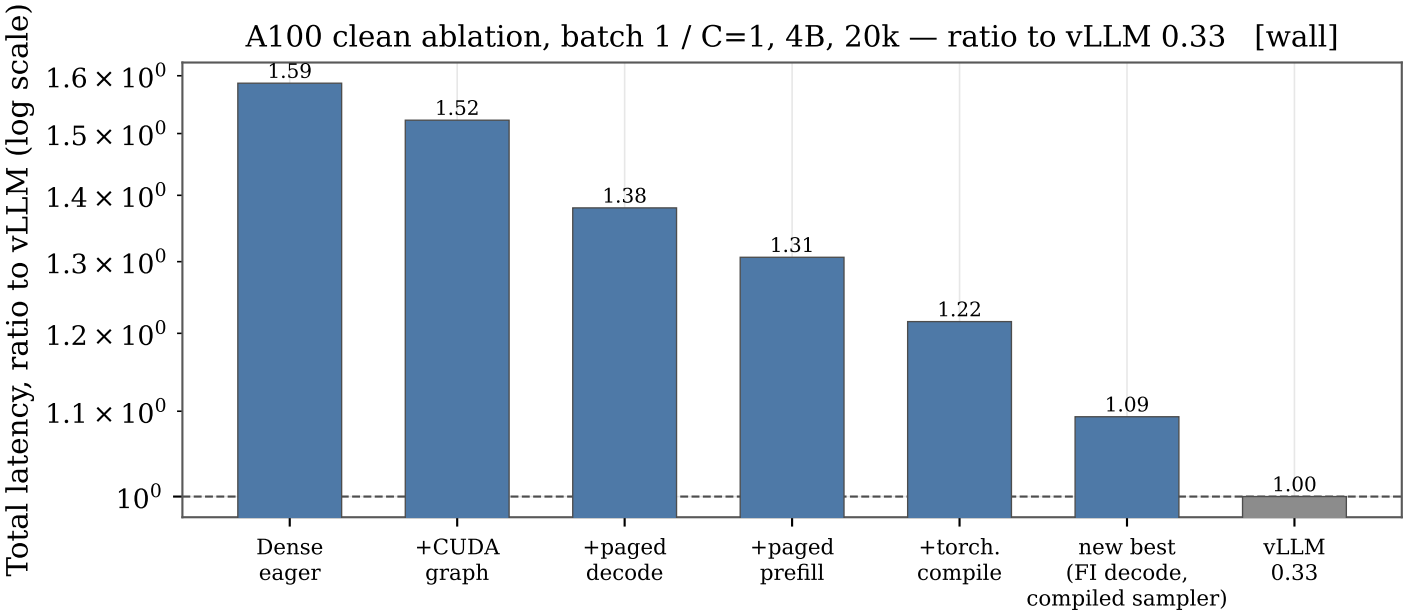
A100 clean ablation, batch 1 / C=1, 4B, 20k — ratio to vLLM 0.33 [wall]
Total latency, ratio to vLLM (log scale)
1.6×10^0
1.5×10^0
1.4×10^0
1.3×10^0
1.2×10^0
1.1×10^0
10^0
1.59
1.52
1.38
1.31
1.22
1.09
1.00
Dense eager
+CUDA graph
+paged decode
+paged prefill
+torch. compile
new best (FI decode, compiled sampler)
vLLM 0.33