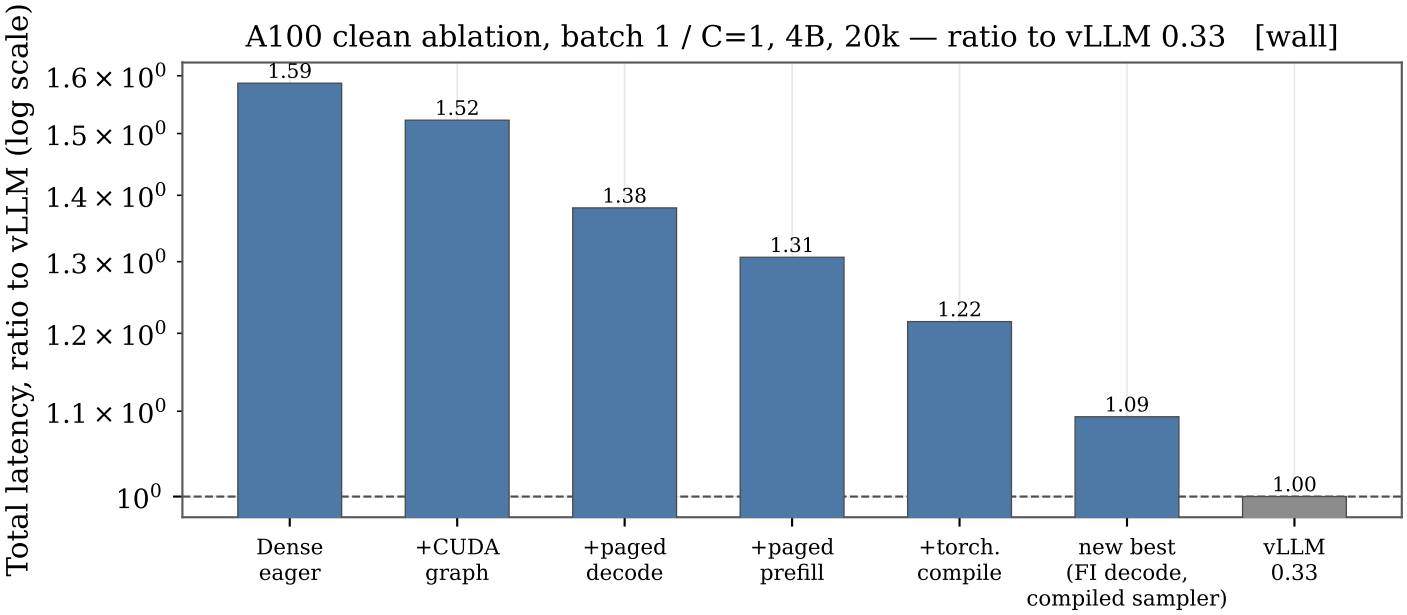
A100 clean ablation, batch 1 / C=1, 4B, 20k — ratio to vLLM 0.33 [wall]
Total latency, ratio to vLLM (log scale)
$1.6 \times 10^0$
$1.5 \times 10^0$
$1.4 \times 10^0$
$1.3 \times 10^0$
$1.2 \times 10^0$
$1.1 \times 10^0$
$10^0$
1.59
1.52
1.38
1.31
1.22
1.09
1.00
Dense eager
+CUDA graph
+paged decode
+paged prefill
+torch. compile
new best (FI decode, compiled sampler)
vLLM 0.33